



THE UNIVERSITY  
of ADELAIDE



CRICOS PROVIDER 00123M

# COMP SCI 7327 Concepts in Artificial Intelligence & Machine Learning -NLP

By Dr Wei Zhang

[adelaide.edu.au](http://adelaide.edu.au)

*seek* LIGHT

# Outline

- Natural Language Processing
- Basic pre-processing
- Natural Language Understanding (NLU)
  - Context-free Grammars

# Natural Language Processing

*“Natural language processing is a range of computational techniques for analyzing and representing naturally occurring texts at one or more levels of linguistic analysis for the purpose of achieving human-like languages processing for a range of particular tasks or applications.”*  
by Liddy (1998)

- Some other names:
  - Computational Linguistics
  - Natural Language Engineering
  - Speech and Text Processing


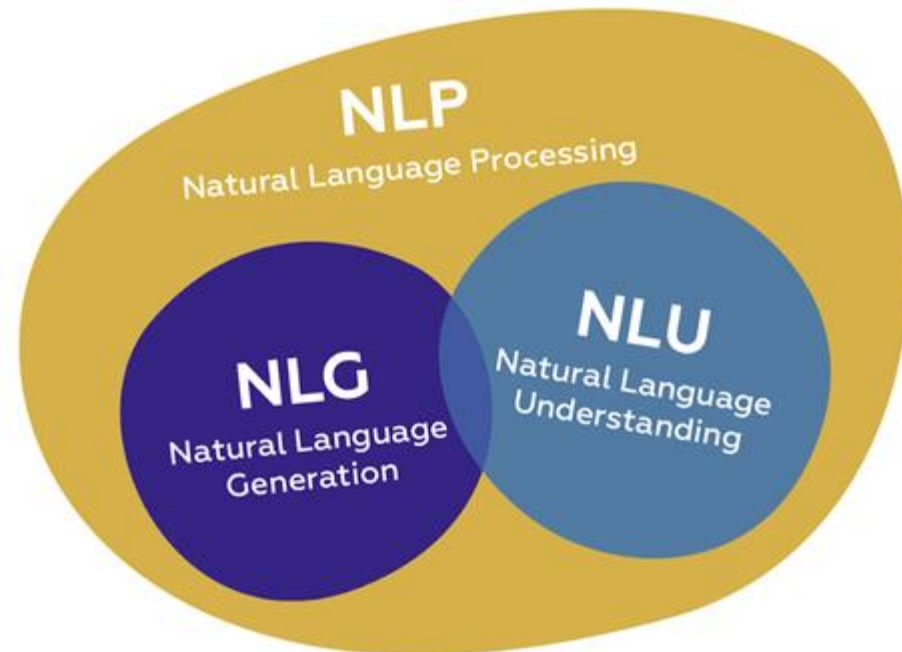


Image source: [algorithmxlab.com](http://algorithmxlab.com)

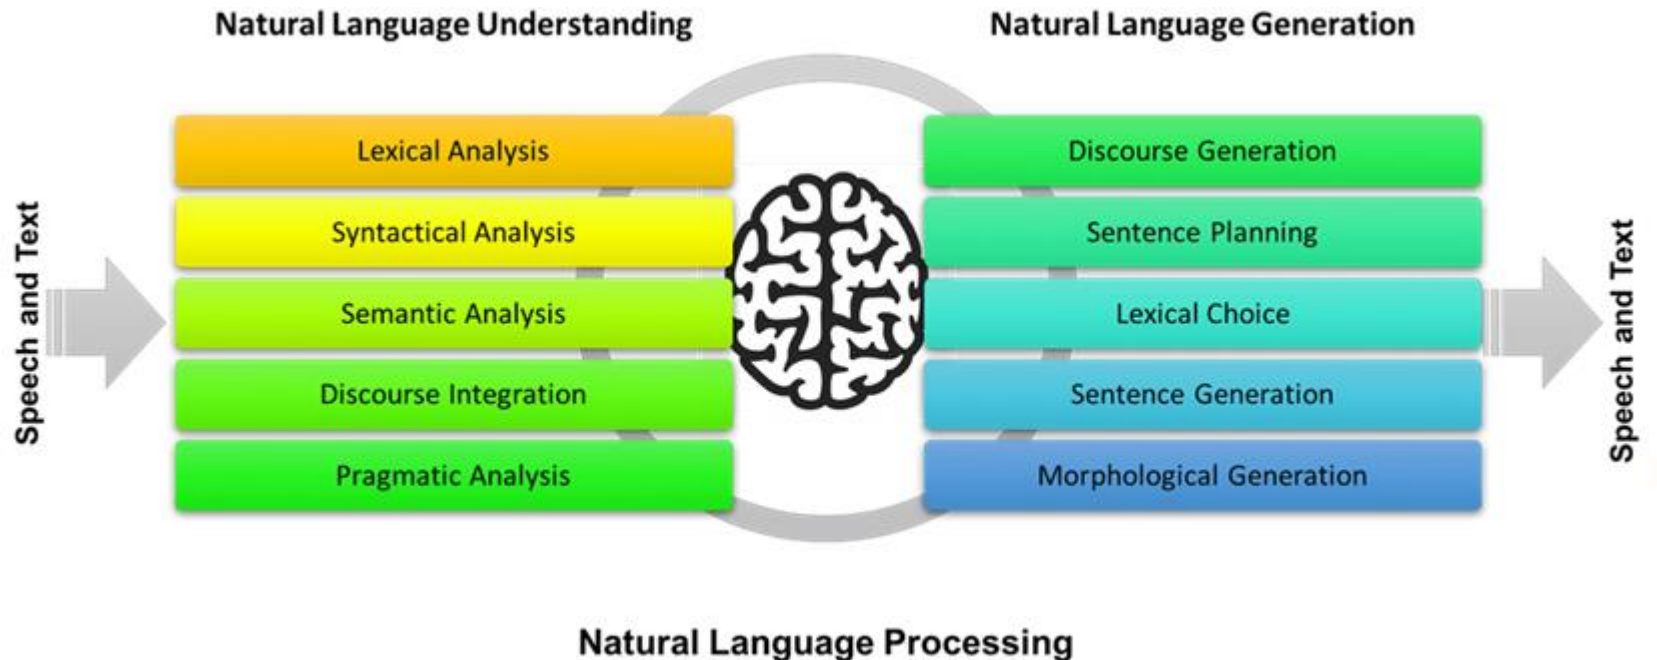


# Natural Language Processing

- NLU
  - Syntactic parsing
  - Coreference resolution
  - Semantic parsing
  - Part-of-speech tagging (POS)
  - Named entity recognition (NER)
  - Natural language inference
  - Relation extraction
  - Text categorization
  - Sentiment analysis
  - ....
- NLG & NLG
  - Paraphrase
  - Dialogue agents
  - Question answering
  - Text summarization
  - Machine translation
  - ...



# Natural Language Processing



- Dialogue system:
  - A typical application to combine NLU and NLG

# Outline

- Natural Language Processing
- Basic pre-processing
- Natural Language Understanding (NLU)
  - Context-free Grammars

# Basic Pre-Processing

- Regular expression
- Word tokenization

# Regular Expressions

- A formal language for specifying text strings
- How can we search for any of these?
  - woodchuck
  - woodchucks
  - Woodchuck
  - Woodchucks





# Regular Expressions: Disjunctions

- Letters inside square brackets []

Pattern	Matches
<code>[wW]oodchuck</code>	Woodchuck, woodchuck
<code>[1234567890]</code>	Any digit

- Ranges `[A-Z]`

Pattern	Matches	Example Patterns Matched
<code>[A-Z]</code>	An upper case letter	<u>D</u> renched Blossoms
<code>[a-z]</code>	A lower case letter	<u>m</u> y beans were impatient
<code>[0-9]</code>	A single digit	Chapter <u>1</u> : Down the Rabbit Hole

# Regular Expressions: In Python

```
import re

words = ['Woodchuck', 'woodchuck', 'chuck', 'wWoodchuck']

pattern = re.compile(r"[wW]oodchuck")

print ("====match/search====")
for wd in words:
    t=re.match(pattern,wd)
    if (t):
        print(t.group(0))
    else:
        print ("Not match \"%s\""% wd)
```

```
====match/search====
Woodchuck
woodchuck
Not match "chuck"
Not match "wWoodchuck"
```

```
strs = ['Drenched Blossoms', 'my beans were impatient', 'Chapter 1: Down the Rabbit Hole']
```

```
pattern = re.compile(r"[A-Z]")
for s in strs:
    t=re.match(pattern,s)
    if (t):
        print(t.group(0))
    else:
        print ("Not match \"%s\""% s)
```

```
D
Not match "my beans were impatient"
C
```

# Regular Expressions: Negation in Disjunction

- Negations `[^Ss]`
  - Carat means negation only when first in []

Pattern	Matches (single character)	Example Patterns Matched
<code>[^A-Z]</code>	Not an upper case letter	O <u>y</u> fn pripetchik
<code>[^Ss]</code>	Neither 'S' nor 's'	<u>I</u> have no exquisite reason
<code>[^e^]</code>	Neither e nor ^	<u>L</u> ook here
<code>a^b</code>	The pattern a carat b	Look up <u>a^b</u> now

# Regular Expressions: More Disjunction


- Woodchucks is another name for groundhog!
- The pipe | for disjunction

Pattern	Matches
<code>groundhog woodchuck</code>	groundhog woodchuck
<code>yours mine</code>	yours mine
<code>a b c</code>	= <code>[abc]</code>
<code>[gG]roundhog [Ww]oodchuck</code>	groundhog Groundhog Woodchuck woodchuck



# Regular Expressions: ? \* + .

Pattern	Matches	Example Patterns Matched
<code>colou?r</code>	Optional previous char	<u>color</u> <u>colour</u>
<code>oo*h!</code>	0 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u> ...
<code>o+h!</code>	1 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u> ...
<code>baa+</code>	=baaa*	<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
<code>beg.n</code>	any character between <i>beg</i> and <i>n</i>	<u>begin</u> <u>beg'n</u> <u>begun</u> <u>beg3n</u>



Stephen C Kleene  
Kleene \*, Kleene +



# Regular Expressions: Anchors <sup>^</sup> <sup>\$</sup>

- The caret <sup>^</sup> matches the start of a line.
- The dollar sign <sup>\$</sup> matches the end of a line.

Pattern	Matches	Example Patterns Matched
<sup>^</sup> [A-Z]	Starting symbol is an upper case letter	<u>P</u> alo Alto
<sup>^</sup> [ ^A-Za-z ]	Starting character is not letter.	<u>1</u> <u>"Hello"</u>
\. <sup>\$</sup>		The end <u>.</u>
. <sup>\$</sup>		The end <u>?    The end!</u>

# Example

- Find me all instances of the word “the” in a text.

the

Misses capitalized examples: *The*

[tT]he

Incorrectly returns *other* or *theology*

[^a-zA-Z]\*[tT]he[^a-zA-Z]|^[tT]he\$

# Example in Python

```
strs = ['It is the cat, that is the dog', 'the dog', 'The dog', 'other thing', 'what is theology', 'the' ]  
  
pattern = re.compile(r"^[a-zA-Z]*[tT]he^[a-zA-Z]*")  
for s in strs:  
    t=re.findall(pattern,s)  
    if (t):  
        print(t)  
    else:  
        print ("Not match \"%s\" % s)
```

```
[' the ', ' the ']  
['the ']  
['The ']  
['the']  
[' the']  
['the']
```

# Regular Expressions

- Regular expressions play a surprisingly large role
  - Sophisticated sequences of regular expressions are often the first model for any text processing text
- For many hard tasks, we use machine learning classifiers
  - But regular expressions are used as features in the classifiers
  - Can be very useful in capturing generalizations

# Tokenization

- How many words?

	Tokens = $N$	Distinct words = $ V $
Switchboard phone conversations	2.4 million	20 thousand
Shakespeare	884,000	31 thousand
Google N-grams	1 trillion	13 million

$N$  = number of tokens

$V$  = vocabulary = number of distinct words

$|V|$  is the size of the vocabulary



# Simple Tokenization in UNIX

- Given a text file, output the word tokens and their frequencies

```
tr -sc 'A-Za-z' '\n' < shakes.txt  
    | sort  
    | uniq -c
```

Change all non-alpha to newlines

Sort in alphabetical order

Merge and count each type

```
1945 A  
72 AARON  
19 ABBESS  
25 Aaron  
6 Abate  
1 Abates  
5 Abbess  
6 Abbey  
3 Abbot  
....
```

# Simple Tokenization in Python

```
import nltk
```

```
sent = "This sentence is going to be tokenized."
```

```
sent.split()
```

```
['This', 'sentence', 'is', 'going', 'to', 'be', 'tokenized.']
```

```
from nltk import word_tokenize  
word_tokenize(sent)
```

```
['This', 'sentence', 'is', 'going', 'to', 'be', 'tokenized', '.']
```

# Issues in Tokenization

- Finland's capital → Finland Finlands Finland's ?
- what're, I'm, isn't → What are, I am, is not
- Hewlett-Packard → Hewlett Packard ?
- state-of-the-art → state of the art ?
- Lowercase → lower-case lowercase lower case ?
- San Francisco → one token or two?
- m.p.h., PhD. → ??

# Tokenization with Regular Expressions

```
import nltk
```

```
text = 'That U.S.A. poster-print costs $12.40...'
```

```
pattern = r'''(?x)          # set flag to allow verbose regexps
    (?:[A-Z]\.)+          # abbreviations, e.g. U.S.A.
    | \w+(?:-\w+)*        # words with optional internal hyphens
    | \$?\d+(?:\.\d+)?%?   # currency and percentages, e.g. $12.40, 82%
    | \.\.\.             # ellipsis
    | [][.,;"'()?:_`-]    # these are separate tokens; includes ], [
    ...
```

```
nltk.regexp_tokenize(text, pattern)
```

```
['That', 'U.S.A.', 'poster-print', 'costs', '$12.40', '...']
```

# Tokenization: language issues

- Chinese and Japanese no spaces between words:
  - Sharapova now lives in US southeastern Florida  
莎拉波娃现在居住在美国东南部的佛罗里达。
  - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
  - Sharapova now lives in US southeastern Florida



# Word Tokenization in Chinese

- Also called **Word Segmentation**
- Chinese words are composed of characters
  - Characters are generally 1 syllable and 1 morpheme (a single unit of meaning ).
  - Average word is 2.4 characters long.
- Standard baseline segmentation algorithm:
  - Maximum Matching

# Maximum Matching Word Segmentation Algorithm

- Given a wordlist of Chinese, and a string.
  - 1) Start a pointer at the beginning of the string
  - 2) Find the longest word in dictionary that matches the string starting at pointer
  - 3) Move the pointer over the word in string
  - 4) Go to 2

# Max-match Segmentation Illustration

- Thecatinthehat the cat in the hat
- Thetabledownthere the table down there
- Doesn't generally work in English! theta bled own there
- But works astonishingly well in Chinese
  - 莎拉波娃现在居住在美国东南部的佛罗里达。
  - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
  - Sharapova now lives in US southeastern Florida
- Modern probabilistic segmentation algorithms even better

# Outline

- Natural Language Processing
- Basic pre-processing
- Natural Language Understanding (NLU)
  - Context-free Grammars

# Natural Language Understanding



Image source: FinTechXpert

Xpert

- Syntactical Analysis
  - The word syntax comes from the Greek *syntaxis*, meaning “setting out together or arrangement”, and refers to the way words are arranged together.
  - Syntactic parsing is the task of recognizing a sentence and assigning a syntactic structure to it.
  - Context free grammars are used for syntactic parsing. And they are the backbone of many formal models of the syntax of natural language.



# Context-free Grammars

- A context-free grammar (CFG) is a set of recursive rewriting rules (or productions) used to generate patterns of strings.
- A CFG consists of the following components:
  - a set of terminal symbols
  - a set of non-terminal symbols
  - a set of productions
  - a start symbol.

# An Example

Context free grammar:

$S \rightarrow NP VP$

$NP \rightarrow Det N \mid Det N PP$

$VP \rightarrow V NP \mid V NP PP$

$PP \rightarrow P NP$


$Det \rightarrow "the" \mid "a"$

$N \rightarrow "man" \mid "dog" \mid "park"$

$V \rightarrow "saw"$


$P \rightarrow "in"$

Syntax Trees:



The sentence to be parsed:

The dog saw a man in the park.



# Syntactic Categories

<b>Symbol</b>	<b>Meaning</b>	<b>Example</b>
S	sentence	<i>the dog saw a man in the park</i>
NP	noun phrase	<i>the dog</i>
VP	verb phrase	<i>saw a man</i>
PP	prepositional phrase	<i>in the park</i>
Det	determiner	<i>the</i>
N	noun	<i>dog</i>
V	verb	<i>saw</i>
P	preposition	<i>in</i>

# Context-free Grammars

- Start symbol:  $S$
- Production rule:  $\{S \rightarrow NP VP, \dots\}$
- Non-terminal symbols:  $\{S, NP, VP, PP, Det, N, V, P, \dots\}$
- Terminal symbols:  $\{"dog", "man", "a", "saw", "in", \dots\}$

# A Context-Free Grammar in Python

```
from nltk import CFG
cfg = CFG.fromstring("""
S -> NP VP
NP -> Det N | Det N PP
VP -> V NP | V NP PP
PP -> P NP
Det -> "the" | "a"
N -> "man" | "dog" | "park"
V -> "saw"
P -> "in"
""")

print(cfg)
```

# Output

Grammar with 13 productions (start state = S)

S -> NP VP

NP -> Det N

NP -> Det N PP

VP -> V NP

VP -> V NP PP

PP -> P NP

Det -> 'the'

Det -> 'a'

N -> 'man'

N -> 'dog'

N -> 'park'

V -> 'saw'

P -> 'in'

# Recursive Descent Parser

- A recursive descent parser is a top-down parser that breaks a high-level goal into several lower-level sub-goals.
- The top-level goal is to find **S**.
- The **S**  $\rightarrow$  **NP VP** production allows the parser to replace this goal with two sub-goals: find an **NP**, then find a **VP**.
- Each sub-goal can in turn be replaced by sub-sub-goals.
- Eventually, this process leads to sub-goals such as: find the determiner **the**.

# Recursive Descent Parser

- The recursive descent parser builds a parse tree during this process.
- With the initial goal (find an **S**), the **S** root node is created.
- As the process recursively expands its goals using the production rules of the grammar, the parse tree is extended downwards.
- During this process, the parser is often forced to choose between alternative productions.
- When a particular production does not work, then the parser has to backtrack.



# Recursive Descent Parsing

- Initial stage:


S

```
S -> NP VP
NP -> Det N
NP -> Det N PP
VP -> V NP
VP -> V NP PP
PP -> P NP
Det -> 'the'
Det -> 'a'
N -> 'man'
N -> 'dog'
N -> 'park'
V -> 'saw'
P -> 'in'
```

.....  
the dog saw a man in the park

# Recursive Descent Parsing

- Second production:




```
S -> NP VP
NP -> Det N
NP -> Det N PP
VP -> V NP
VP -> V NP PP
PP -> P NP
Det -> 'the'
Det -> 'a'
N -> 'man'
N -> 'dog'
N -> 'park'
V -> 'saw'
P -> 'in'
```

.....  
the dog saw a man in the park

# Recursive Descent Parsing

- Matching determiner *the*:




```
S -> NP VP
NP -> Det N
NP -> Det N PP
VP -> V NP
VP -> V NP PP
PP -> P NP
Det -> 'the'
Det -> 'a'
N -> 'man'
N -> 'dog'
N -> 'park'
V -> 'saw'
P -> 'in'
```

the dog saw a man in the park

# Recursive Descent Parsing


- Cannot match *man*:



```
S -> NP VP
NP -> Det N
NP -> Det N PP
VP -> V NP
VP -> V NP PP
PP -> P NP
Det -> 'the'
Det -> 'a'
N -> 'man'
N -> 'dog'
N -> 'park'
V -> 'saw'
P -> 'in'
```

# Recursive Descent Parsing


- Complete parse:



S -> NP VP  
NP -> Det N  
NP -> Det N PP  
VP -> V NP  
VP -> V NP PP  
PP -> P NP  
Det -> 'the'  
Det -> 'a'  
N -> 'man'  
N -> 'dog'  
N -> 'park'  
V -> 'saw'  
P -> 'in'

# Recursive Descent Parsing


- Backtracking for second interpretation:



S -> NP VP  
NP -> Det N  
NP -> Det N PP  
VP -> V NP  
VP -> V NP PP  
PP -> P NP  
Det -> 'the'  
Det -> 'a'  
N -> 'man'  
N -> 'dog'  
N -> 'park'  
V -> 'saw'  
P -> 'in'

# First Syntax Tree


- Ambiguity: prepositional phrase attachment



S -> NP VP  
NP -> Det N  
NP -> Det N PP  
VP -> V NP  
VP -> V NP PP  
PP -> P NP  
Det -> 'the'  
Det -> 'a'  
N -> 'man'  
N -> 'dog'  
N -> 'park'  
V -> 'saw'  
P -> 'in'

# Second Syntax Tree

- Ambiguity: prepositional phrase attachment



S -> NP VP  
NP -> Det N  
NP -> Det N PP  
VP -> V NP  
VP -> V NP PP  
PP -> P NP  
Det -> 'the'  
Det -> 'a'  
N -> 'man'  
N -> 'dog'  
N -> 'park'  
V -> 'saw'  
P -> 'in'



# A Recursive Descent Parser in Python

```
# Recursive Descent Parser (top-down, depth-first)
from nltk import RecursiveDescentParser

sent = "the dog saw a man in the park".split()

rd_parser = RecursiveDescentParser(cfg)

for tree in rd_parser.parse(sent):
    print(tree)
```

# Output

```
(S
  (NP (Det the) (N dog))
  (VP
    (V saw)
    (NP (Det a) (N man) (PP (P in) (NP (Det the) (N park))))))
```

```
(S
  (NP (Det the) (N dog))
  (VP
    (V saw)
    (NP (Det a) (N man))
    (PP (P in) (NP (Det the) (N park))))))
```

# Probabilistic Context-Free Grammars

- A probabilistic context-free grammar (or *PCFG*) is a context-free grammar that associates a probability with each of its productions.
- It generates the same set of parses for a text that the corresponding context-free grammar does, and assigns a probability to each parse.
- The probability of a parse generated by a PCFG is simply the product of the probabilities of the productions used to generate it.

# A PCFG in Python

```
from nltk import PCFG
pcfg = PCFG.fromstring("""
S -> NP VP      [1.0]
NP -> Det N      [0.7]
NP -> Det N PP   [0.3]
VP -> V NP       [0.6]
VP -> V NP PP    [0.4]
PP -> P NP       [1.0]
Det -> "the"     [0.6]
Det -> "a"       [0.4]
N -> "man"       [0.4]
N -> "dog"       [0.3]
N -> "park"      [0.3]
V -> "saw"       [1.0]
P -> "in"        [1.0]
""")
```

# Viterbi Parser in Python

```
from nltk import ViterbiParser

sent = "the dog saw a man in the park".split()

viterbi_parser = ViterbiParser(pcfg, trace=2)

for tree in viterbi_parser.parse(sent):
    print(tree)
```

# Output

```
(S
  (NP (Det the) (N dog))
  (VP
    (V saw)
    (NP (Det a) (N man))
    (PP (P in) (NP (Det the) (N park)))) (p=0.000711245)
```

# Some Online Parsers

- Stanford Parser:

- <http://nlp.stanford.edu:8080/parser/index.jsp>

```
(ROOT
  (S
    (NP (DT the) (NN dog))
    (VP (VBD saw)
      (NP (DT a) (NN man))
      (PP (IN in)
        (NP (DT the) (NN park))))))
```

- <http://www.link.cs.cmu.edu/link/submit-sentence-4.html>

```
(S (NP the dog)
  (VP saw
    (NP a man)
    (PP in
      (NP the park))))
```

---

# References